

## task-3

February 6, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv('Churn_Modelling.csv')
```

```
[3]: df.head()
```

```
[3]:   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0          1    15634602  Hargrave         619      France  Female   42
1          2    15647311    Hill         608      Spain  Female   41
2          3    15619304    Onio         502      France  Female   42
3          4    15701354    Boni         699      France  Female   39
4          5    15737888  Mitchell         850      Spain  Female   43
```

```
   Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0        2     0.00              1          1              1
1        1  83807.86              1          0              1
2        8 159660.80              3          1              0
3        1     0.00              2          0              0
4        2 125510.82              1          1              1
```

```
   EstimatedSalary  Exited
0      101348.88      1
1      112542.58      0
2      113931.57      1
3       93826.63      0
4       79084.10      0
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column              Non-Null Count  Dtype
---  -
0   RowNumber           10000 non-null  int64
```

```

1  CustomerId      10000 non-null  int64
2  Surname         10000 non-null  object
3  CreditScore     10000 non-null  int64
4  Geography       10000 non-null  object
5  Gender          10000 non-null  object
6  Age            10000 non-null  int64
7  Tenure          10000 non-null  int64
8  Balance         10000 non-null  float64
9  NumOfProducts  10000 non-null  int64
10 HasCrCard       10000 non-null  int64
11 IsActiveMember 10000 non-null  int64
12 EstimatedSalary 10000 non-null  float64
13 Exited          10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```


```
[5]: df.isnull().sum()
```

```

[5]: RowNumber      0
    CustomerId      0
    Surname         0
    CreditScore     0
    Geography       0
    Gender          0
    Age            0
    Tenure          0
    Balance         0
    NumOfProducts  0
    HasCrCard       0
    IsActiveMember  0
    EstimatedSalary 0
    Exited          0
    dtype: int64

```

```
[6]: df.hist(bins=5,figsize=(15,15), color='green')
plt.show()
```



```
[7]: df.columns
```

```
[7]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
          'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
          'IsActiveMember', 'EstimatedSalary', 'Exited'],
          dtype='object')
```

```
[8]: df.head()
```

```
[8]:   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0         1    15634602  Hargrave         619    France  Female   42
1         2    15647311    Hill         608    Spain  Female   41
2         3    15619304    Onio         502    France  Female   42
```

3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
[9]: df.drop(columns=['RowNumber', 'CustomerId', 'Surname'], axis=1, inplace=True)
```

```
[10]: df.head()
```

```
[10]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	Female	42	2	0.00	1	
1	608	Spain	Female	41	1	83807.86	1	
2	502	France	Female	42	8	159660.80	3	
3	699	France	Female	39	1	0.00	2	
4	850	Spain	Female	43	2	125510.82	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0

```
[11]: from sklearn.preprocessing import LabelEncoder
lr=LabelEncoder()
df['Geography']=lr.fit_transform(df['Geography'])
df['Gender']=lr.fit_transform(df['Gender'])
```

```
[12]: df.head()
```

```
[12]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	0	0	42	2	0.00	1	
1	608	2	0	41	1	83807.86	1	
2	502	0	0	42	8	159660.80	3	
3	699	0	0	39	1	0.00	2	

4	850	2	0	43	2	125510.82	1
---	-----	---	---	----	---	-----------	---

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0

```
[13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CreditScore            10000 non-null  int64
1   Geography              10000 non-null  int32
2   Gender                 10000 non-null  int32
3   Age                   10000 non-null  int64
4   Tenure                 10000 non-null  int64
5   Balance                10000 non-null  float64
6   NumOfProducts          10000 non-null  int64
7   HasCrCard              10000 non-null  int64
8   IsActiveMember         10000 non-null  int64
9   EstimatedSalary        10000 non-null  float64
10  Exited                  10000 non-null  int64
dtypes: float64(2), int32(2), int64(7)
memory usage: 781.4 KB
```

```
[14]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, classification_report, \
      ↪confusion_matrix
```

```
[15]: x = df.drop('Exited', axis=1)
      y = df['Exited']
```

```
[16]: x.head()
```

```
[16]:   CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  \
0         619          0      0    42      2      0.00           1
1         608          2      0    41      1  83807.86           1
2         502          0      0    42      8  159660.80           3
3         699          0      0    39      1      0.00           2
4         850          2      0    43      2  125510.82           1
```

	HasCrCard	IsActiveMember	EstimatedSalary
0	1	1	101348.88
1	0	1	112542.58
2	1	0	113931.57
3	0	0	93826.63
4	1	1	79084.10

```
[17]: y.head()
```

```
[17]: 0    1
      1    0
      2    1
      3    0
      4    0
      Name: Exited, dtype: int64
```

```
[18]: x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2,
      ↪random_state=123)
      print('x_train', x_train.shape)
      print('x_test', x_test.shape)
      print('y_train', y_train.shape)
      print('y_test', y_test.shape)
```

```
x_train (8000, 10)
x_test (2000, 10)
y_train (8000,)
y_test (2000,)
```

```
[19]: lr= LogisticRegression()
      lr.fit(x_train,y_train)
      y_pred = lr.predict(x_test)
      accuracy = accuracy_score(y_test,y_pred)
      class_report = classification_report(y_test,y_pred)
      conf_matrix = confusion_matrix(y_test,y_pred)
```

```
[20]: print(f'The accuracy score of churn prediction is : {accuracy : .2f}%\n ')
      print(f'The classification report of churn prediction is \n: {class_report}')
      print(f'The confusion matrix of churn prediction is:\n{conf_matrix}')
```

The accuracy score of churn prediction is : 0.78%

The classification report of churn prediction is


:		precision	recall	f1-score	support
	0	0.80	0.98	0.88	1586
	1	0.31	0.04	0.07	414

accuracy			0.78	2000
macro avg	0.56	0.51	0.47	2000
weighted avg	0.70	0.78	0.71	2000

The confusion matrix of churn prediction is:

```
[[1549  37]
 [ 397  17]]
```

```
[21]: sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g')
plt.show()
```



```
[ ]: 
```

```
[ ]: 
```

```
[ ]: 
```